

banknote_forgery

August 30, 2017

In this notebook, we will build a neural net capable of discerning real banknotes from forged ones based some of their statistical properties

We first import the necessary modules

```
In [145]: import pandas as pd
import numpy as np
from IPython.display import display
import matplotlib.pyplot as plt
from sklearn.metrics import f1_score
from sklearn.model_selection import train_test_split
```

We load the dataset as a pandas dictionary and extract the features and target as separate datasets

```
In [21]: data=pd.read_csv("data.csv")
features=data.drop('forge', axis=1)
target=data['forge']
```


1 Preprocessing and Viewing the Data

Each element of the dataset contains 4 features which describe the statistical properties of a banknote after scanning it and applying a wavelet transform. We begin by exhibiting some descriptive statistical properties of the dataset:

	var	skew	curt	ent
count	1372.000000	1372.000000	1372.000000	1372.000000
mean	0.433735	1.922353	1.397627	-1.191657
std	2.842763	5.869047	4.310030	2.101013
min	-7.042100	-13.773100	-5.286100	-8.548200
25%	-1.773000	-1.708200	-1.574975	-2.413450
50%	0.496180	2.319650	0.616630	-0.586650
75%	2.821475	6.814625	3.179250	0.394810
max	6.824800	12.951600	17.927400	2.449500

In []: To visualize these descriptive statistics, we display the boxplot:

```
In [23]: features.boxplot()
plt.show()
```



It is clear that datapoints with a kurtosis of 10 or higher are outliers. We compute these

```
In [24]: features[(features['curt']>10)]
```

```
Out [24]:
```

	var	skew	curt	ent
765	-3.8483	-12.8047	15.6824	-1.281000
766	-3.5681	-8.2130	10.0830	0.967650
780	-3.5801	-12.9309	13.1779	-2.567700
815	-3.1128	-6.8410	10.7402	-1.017200
816	-4.8554	-5.9037	10.9818	-0.821990
820	-4.0025	-13.4979	17.6772	-3.320200
821	-4.0173	-8.3123	12.4547	-1.437500
826	-4.2110	-12.4736	14.9704	-1.388400
827	-3.8073	-8.0971	10.1772	0.650840
841	-3.8858	-12.8461	12.7957	-3.135300
876	-3.5916	-6.2285	10.2389	-1.154300
877	-5.1216	-5.3118	10.3846	-1.061200
881	-4.4861	-13.2889	17.3087	-3.219400
882	-4.3876	-7.7267	11.9655	-1.454300
887	-3.2692	-12.7406	15.5573	-0.141820
902	-2.8957	-12.0205	11.9149	-2.755200
937	-2.9020	-7.6563	11.8318	-0.842680
938	-4.3773	-5.5167	10.9390	-0.408200
942	-3.3793	-13.7731	17.9274	-2.032300
943	-3.1273	-7.1121	11.3897	-0.083634
948	-3.4917	-12.1736	14.3689	-0.616390
949	-3.1158	-8.6289	10.4403	0.971530

963	-3.3863	-12.9889	13.0545	-2.720200
998	-3.0866	-6.6362	10.5405	-0.891820
999	-4.7331	-6.1789	11.3880	-1.074100
1003	-3.8203	-13.0551	16.9583	-2.305200
1004	-3.7181	-8.5089	12.3630	-0.955180
1009	-3.5713	-12.4922	14.8881	-0.470270
1023	-1.7713	-10.7665	10.2184	-1.004300
1024	-3.0061	-12.2377	11.9552	-2.160300
...
1121	-4.6765	-5.6636	10.9690	-0.334490
1125	-3.5985	-13.6593	17.6052	-2.492700
1126	-3.3582	-7.2404	11.4419	-0.571130
1131	-4.0214	-12.8006	15.6199	-0.956470
1132	-3.3884	-8.2150	10.3315	0.981870
1146	-3.7300	-12.9723	12.9817	-2.684000
1181	-3.5895	-6.5720	10.5251	-0.163810
1182	-5.0477	-5.8023	11.2440	-0.390100
1186	-4.2440	-13.0634	17.1116	-2.801700
1187	-4.0218	-8.3040	12.5550	-1.509900
1192	-4.4018	-12.9371	15.6559	-1.680600
1193	-3.7573	-8.2916	10.3032	0.380590
1207	-3.7930	-12.7095	12.7957	-2.825000
1242	-3.6053	-5.9740	10.0916	-0.828460
1243	-5.0676	-5.1877	10.4266	-0.867250
1247	-4.4775	-13.0303	17.0834	-3.034500
1248	-4.1958	-8.1819	12.1291	-1.601700
1253	-4.5531	-12.5854	15.4417	-1.498300
1268	-3.9411	-12.8792	13.0597	-3.312500
1303	-3.9297	-6.0816	10.0958	-1.014700
1304	-5.2943	-5.1463	10.3332	-1.118100
1308	-4.6338	-12.7509	16.7166	-3.216800
1309	-4.2887	-7.8633	11.8387	-1.897800
1314	-3.5060	-12.5667	15.1606	-0.752160
1315	-2.9498	-8.2730	10.2646	1.162900
1329	-2.9672	-13.2869	13.4727	-2.627100
1364	-2.8391	-6.6300	10.4849	-0.421130
1365	-4.5046	-5.8126	10.8867	-0.528460
1369	-3.7503	-13.4586	17.5932	-2.777100
1370	-3.5637	-8.3827	12.3930	-1.282300

[67 rows x 4 columns]

It turns out there are 67 outliers with a curtosis higher than 10. We remove them from the dataset.


```
In [248]: new_data=data[(features['curt']<10)]
          new_features=new_data.drop(['forge'], axis=1)
          new_target=new_data['forge']
```

Since we intend to build a learner using a neural net, we first compute the various correlations between the data, as correlated features could result in overfitting:

```
In [256]: from pandas.tools.plotting import scatter_matrix
```

```
correlation=new_features.corr(method='pearson')
print(correlation)
scatter_matrix(new_features)
plt.show()
```

	var	skew	curt	ent
var	1.000000	0.137317	-0.239365	0.292094
skew	0.137317	1.000000	-0.722368	-0.613270
curt	-0.239365	-0.722368	1.000000	0.427658
ent	0.292094	-0.613270	0.427658	1.000000



It seems there is a high correlation between skewness and kurtosis, which is to be expected. The other features seem to be uncorrelated

to visualize the features, we first perform a principal component analysis to reduce the 4-dimensional feature space to 3, and then plot the transformed features.


```
In [253]: from sklearn.decomposition import PCA
          from mpl_toolkits.mplot3d import Axes3D
```

```
pca=PCA(n_components=3)
pca.fit(new_features)
first_pc=pca.components_[0]
second_pc=pca.components_[1]
third_pc=pca.components_[2]
reduced_data=pca.transform(new_features)

fig = plt.figure(figsize=(10,10))
ax = Axes3D(fig)

x=reduced_data[:,0]
y=reduced_data[:,1]
z=reduced_data[:,2]

ax.scatter(x,y,z)
plt.show()
```



2 Building the Neural Net

We first randomize the data and split into test and training set 2/3-1/3

```
In [29]: X_train, X_test, Y_train, Y_test = train_test_split(new_features, new_target,
```

Next, we import the necessary functionality from keras:

```
In [135]: from keras.utils import np_utils
          from keras.callbacks import ModelCheckpoint
          from keras.models import Sequential
          from keras.layers import Dense, Dropout
          from keras import optimizers
```

We reshape the training and testing data in a manner that is compatible to keras

```
In [257]: ##translate from dataframe to numpy array

x_train=X_train.as_matrix()
x_test=X_test.as_matrix()
y_train = np_utils.to_categorical(Y_train, 2)
y_test = np_utils.to_categorical(Y_test, 2)
```

Next, to view the results later on, we define a function that plots the learning curves for both the loss and accuracy

```
In [258]: def learning_curves(history):
    plt.plot(history.history['acc'])
    plt.plot(history.history['val_acc'])
    plt.title('model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()
    # summarize history for loss
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()
```

Next, we define a function that will compile and fit a neural network according to a pre-described optimizer

```
In [259]: def comp_fit(model,sgd):
    epochs=10
    model.compile(optimizer=sgd, loss='categorical_crossentropy', metrics=
    checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.fr
    history=model.fit(x_train, y_train, validation_data=(x_test, y_test),
    return history
```



As a first simple model, we implement logistic regression according using the standard rmsprop optimizer

```
In [266]: ##build a logistical regression first
logmodel=Sequential()
logmodel.add(Dense(2,input_dim=4,activation='sigmoid'))
history=comp_fit(logmodel, 'rmsprop')
learning_curves(history)
```

```

Train on 874 samples, validate on 431 samples
Epoch 1/10
660/874 [=====>...] - ETA: 0s - loss: 1.2631 - acc: 0.1712Epoch 0000
874/874 [=====] - 3s - loss: 1.2498 - acc: 0.1682 - val_loss: 1.2498
Epoch 2/10
620/874 [=====>...] - ETA: 0s - loss: 1.2031 - acc: 0.1661Epoch 0000
874/874 [=====] - 0s - loss: 1.1584 - acc: 0.1716 - val_loss: 1.1584
Epoch 3/10
600/874 [=====>...] - ETA: 0s - loss: 1.1072 - acc: 0.2000Epoch 0000
874/874 [=====] - 0s - loss: 1.0784 - acc: 0.2002 - val_loss: 1.0784
Epoch 4/10
560/874 [=====>...] - ETA: 0s - loss: 1.0393 - acc: 0.2143Epoch 0000
874/874 [=====] - 0s - loss: 1.0064 - acc: 0.2151 - val_loss: 1.0064
Epoch 5/10
540/874 [=====>...] - ETA: 0s - loss: 0.9419 - acc: 0.2630Epoch 0000
874/874 [=====] - 0s - loss: 0.9371 - acc: 0.2712 - val_loss: 0.9371
Epoch 6/10
780/874 [=====>...] - ETA: 0s - loss: 0.8876 - acc: 0.3128Epoch 0000
874/874 [=====] - 0s - loss: 0.8741 - acc: 0.3146 - val_loss: 0.8741
Epoch 7/10
600/874 [=====>...] - ETA: 0s - loss: 0.8375 - acc: 0.3883Epoch 0000
874/874 [=====] - 0s - loss: 0.8147 - acc: 0.4211 - val_loss: 0.8147
Epoch 8/10
620/874 [=====>...] - ETA: 0s - loss: 0.7688 - acc: 0.5097Epoch 0000
874/874 [=====] - 0s - loss: 0.7602 - acc: 0.5286 - val_loss: 0.7602
Epoch 9/10
740/874 [=====>...] - ETA: 0s - loss: 0.7125 - acc: 0.6270Epoch 0000
874/874 [=====] - 0s - loss: 0.7083 - acc: 0.6362 - val_loss: 0.7083
Epoch 10/10
580/874 [=====>...] - ETA: 0s - loss: 0.6625 - acc: 0.7172Epoch 0000
874/874 [=====] - 0s - loss: 0.6552 - acc: 0.7094 - val_loss: 0.6552

```

The model seems to fit the data rather nicely with an accuracy of 70%, and a loss function that decays nicely to 0.7. To improve the accuracy, we change the optimizer to an Adam optimizer

```
In [272]: log_model=Sequential()  
          log_model.add(Dense(2,input_dim=4,activation='sigmoid'))  
          history=comp_fit(log_model,'Adam')  
          learning_curves(history)
```

Train on 874 samples, validate on 431 samples

Epoch 1/10

840/874 [=====>...] - ETA: 0s - loss: 1.2969 - acc: 0.5560Epoch

874/874 [=====] - 4s - loss: 1.3047 - acc: 0.5526 - val_lo

Epoch 2/10

820/874 [=====>...] - ETA: 0s - loss: 1.1068 - acc: 0.5744Epoch

874/874 [=====] - 0s - loss: 1.0836 - acc: 0.5767 - val_lo

Epoch 3/10

780/874 [=====>...] - ETA: 0s - loss: 0.8938 - acc: 0.5910Epoch

874/874 [=====] - 0s - loss: 0.8889 - acc: 0.5973 - val_lo

Epoch 4/10

780/874 [=====>...] - ETA: 0s - loss: 0.7539 - acc: 0.6372Epoch

874/874 [=====] - 0s - loss: 0.7283 - acc: 0.6522 - val_lo

Epoch 5/10

820/874 [=====>...] - ETA: 0s - loss: 0.5914 - acc: 0.6963Epoch

874/874 [=====] - 0s - loss: 0.5932 - acc: 0.6934 - val_lo

Epoch 6/10

840/874 [=====>...] - ETA: 0s - loss: 0.4947 - acc: 0.7298Epoch

874/874 [=====] - 0s - loss: 0.4916 - acc: 0.7300 - val_lo

Epoch 7/10

700/874 [=====>...] - ETA: 0s - loss: 0.4148 - acc: 0.7700Epoch

874/874 [=====] - 0s - loss: 0.4206 - acc: 0.7654 - val_lo

Epoch 8/10

760/874 [=====>...] - ETA: 0s - loss: 0.3852 - acc: 0.8158Epoch

874/874 [=====] - 0s - loss: 0.3782 - acc: 0.8238 - val_lo

Epoch 9/10



800/874 [=====>...] - ETA: 0s - loss: 0.3523 - acc: 0.8400Epoch

874/874 [=====] - 0s - loss: 0.3523 - acc: 0.8455 - val_lo

Epoch 10/10

860/874 [=====>.] - ETA: 0s - loss: 0.3372 - acc: 0.8558Epoch

874/874 [=====] - 0s - loss: 0.3352 - acc: 0.8570 - val_lo



as we can see, the accuracy has increased to 85%. To further improve the network, we include a hidden layer

```
In [271]: log_model_hl=Sequential()
          log_model_hl.add(Dense(4,input_dim=4,activation='sigmoid'))
          log_model_hl.add(Dense(2,input_dim=4,activation='sigmoid'))
          history=comp_fit(log_model_hl,'Adam')
          learning_curves(history)
```

Train on 874 samples, validate on 431 samples

Epoch 1/10

740/874 [=====>...] - ETA: 0s - loss: 0.6170 - acc: 0.5797Epoch

874/874 [=====] - 4s - loss: 0.6157 - acc: 0.5812 - val_lo

Epoch 2/10

700/874 [=====>...] - ETA: 0s - loss: 0.5987 - acc: 0.6100Epoch 0

874/874 [=====] - 0s - loss: 0.5987 - acc: 0.6076 - val_lo

Epoch 3/10

660/874 [=====>...] - ETA: 0s - loss: 0.5830 - acc: 0.6485Epoch 000

874/874 [=====] - 0s - loss: 0.5817 - acc: 0.6556 - val_lo

Epoch 4/10

700/874 [=====>...] - ETA: 0s - loss: 0.5657 - acc: 0.7171Epoch 0

874/874 [=====] - 0s - loss: 0.5642 - acc: 0.7220 - val_lo

Epoch 5/10

680/874 [=====>...] - ETA: 0s - loss: 0.5512 - acc: 0.7838Epoch 00

874/874 [=====] - 0s - loss: 0.5461 - acc: 0.7860 - val_lo

Epoch 6/10

860/874 [=====>.] - ETA: 0s - loss: 0.5280 - acc: 0.8151Epoch

874/874 [=====] - 0s - loss: 0.5271 - acc: 0.8181 - val_lo

Epoch 7/10

680/874 [=====>...] - ETA: 0s - loss: 0.5097 - acc: 0.8353Epoch 00

874/874 [=====] - 0s - loss: 0.5065 - acc: 0.8387 - val_lo

Epoch 8/10

660/874 [=====>...] - ETA: 0s - loss: 0.4913 - acc: 0.8606Epoch 000

874/874 [=====] - 0s - loss: 0.4845 - acc: 0.8684 - val_lo

Epoch 9/10



840/874 [=====>..] - ETA: 0s - loss: 0.4603 - acc: 0.8857Epoch

874/874 [=====] - 0s - loss: 0.4605 - acc: 0.8867 - val_lo

Epoch 10/10

700/874 [=====>...] - ETA: 0s - loss: 0.4369 - acc: 0.9043Epoch 0

874/874 [=====] - 0s - loss: 0.4346 - acc: 0.9050 - val_lo



The accuracy has now improved to 90%. We finally include a dropout to see how this changes the accuracy:

```
In [201]: log_model_hl_dp=Sequential()  
          log_model_hl_dp.add(Dense(4,input_dim=4,activation='sigmoid'))  
          log_model_hl_dp.add(Dropout(0.5))  
          log_model_hl_dp.add(Dense(2,input_dim=4,activation='sigmoid'))  
          history=comp_fit(log_model_hl_dp, 'Adam')  
          learning_curves(history)
```

Train on 874 samples, validate on 431 samples

Epoch 1/10

860/874 [=====>.] - ETA: 0s - loss: 0.7044 - acc: 0.4721Epoch

874/874 [=====] - 4s - loss: 0.7036 - acc: 0.4760 - val_lo

Epoch 2/10

820/874 [=====>..] - ETA: 0s - loss: 0.6871 - acc: 0.5354Epoch

874/874 [=====] - 0s - loss: 0.6863 - acc: 0.5366 - val_lo

Epoch 3/10

840/874 [=====>..] - ETA: 0s - loss: 0.6704 - acc: 0.5929Epoch

874/874 [=====] - 0s - loss: 0.6699 - acc: 0.5904 - val_lo

Epoch 4/10

860/874 [=====>.] - ETA: 0s - loss: 0.6696 - acc: 0.5767Epoch

874/874 [=====] - 0s - loss: 0.6699 - acc: 0.5744 - val_lo

Epoch 5/10

820/874 [=====>..] - ETA: 0s - loss: 0.6574 - acc: 0.5939Epoch

874/874 [=====] - 0s - loss: 0.6608 - acc: 0.5904 - val_lo

Epoch 6/10

840/874 [=====>..] - ETA: 0s - loss: 0.6457 - acc: 0.6226Epoch

874/874 [=====] - 0s - loss: 0.6461 - acc: 0.6178 - val_lo

Epoch 7/10

800/874 [=====>...] - ETA: 0s - loss: 0.6390 - acc: 0.6138Epoch

874/874 [=====] - 0s - loss: 0.6418 - acc: 0.6076 - val_lo

Epoch 8/10

760/874 [=====>...] - ETA: 0s - loss: 0.6285 - acc: 0.6474Epoch

874/874 [=====] - 0s - loss: 0.6325 - acc: 0.6362 - val_lo

Epoch 9/10



700/874 [=====>...] - ETA: 0s - loss: 0.6207 - acc: 0.6600Epoch

874/874 [=====] - 0s - loss: 0.6213 - acc: 0.6510 - val_lo

Epoch 10/10

800/874 [=====>...] - ETA: 0s - loss: 0.6029 - acc: 0.6837Epoch

874/874 [=====] - 0s - loss: 0.6008 - acc: 0.6865 - val_lo



Not surprisingly, this decreases the accuracy, as the network was already learning almost optimally

we finally create a function that predicts whether a banknote is real depending on the model and check it on both a real and forged banknote:

```
In [310]: def classify(features):
           features=np.array([features])
           prediction=log_model_h1.predict(features)
           prediction=np.argmax(prediction[0])
           if prediction == 1:
               print('the banknote is forged')
           else:
               print('the banknote is real')

In [311]: classify([3.621600,8.66610,-2.807300,-0.446990] )
           classify([-1.747900, -5.82300,5.869900 ,1.212000])

the banknote is real
the banknote is forged
```